



LangGraph, LangSmith & Multi-Agent Architecture

Surendra Panpaliya

Generative AI

Gen-AI

Agenda



Why orchestration matters:



connecting multiple agents and workflows



LangGraph Architecture:

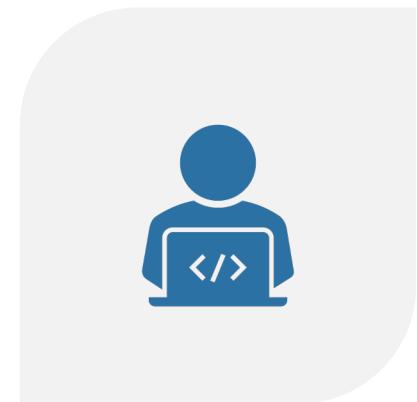


state graph, nodes, and conditional flows

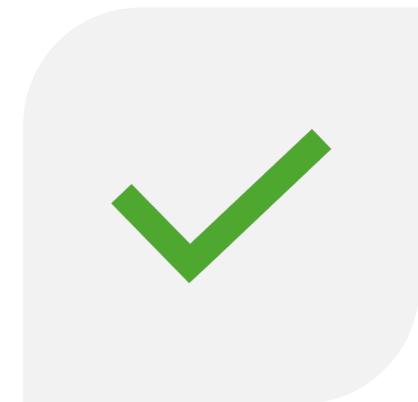
Agenda



LANGSMITH

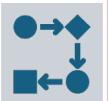


TRACING,
DEBUGGING, AND



OBSERVABILITY OF
LLM PIPELINES

Demo



Visualizing a workflow:



“Data → Retriever → Validator → Summarizer → Notifier”



Implementing decision branches



using LangGraph visual editor

Demo



Multi-Agent Orchestration:



Planner, Reviewer, and Reporter agents (no coding)



Real-time demonstration of AI pipeline monitoring using LangSmith UI



Compliance audit example:



tracking AI outputs and data lineage

Agentic AI & Task Chaining

What is
Agentic AI?

Core
components:

Agents

Tools

Planner

Executor

What is Agentic AI?

Agentic AI = Smart AI agents that

Understand goals

Plan steps

Execute tasks

Work together (like a team)

What is Agentic AI?

New way of using AI

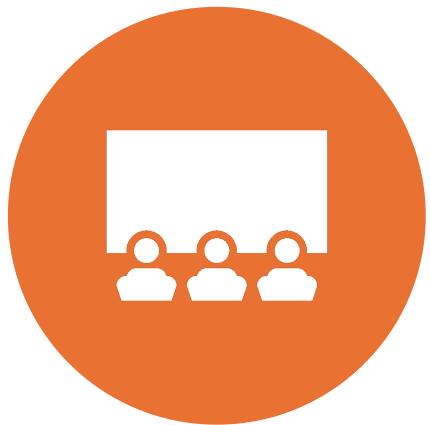
to plan, execute, and manage tasks

independently

like an intelligent assistant

that thinks and acts.

What is Agentic AI?



**ACT LIKE A RESPONSIBLE
ASSISTANT,**



**DOING TASKS
INDEPENDENTLY**



**BASED ON GOALS YOU
GIVE IT.**

Agentic AI uses Agents that can



Understand goals



Break them into tasks



Use tools or APIs to complete tasks



Adjust actions based on results

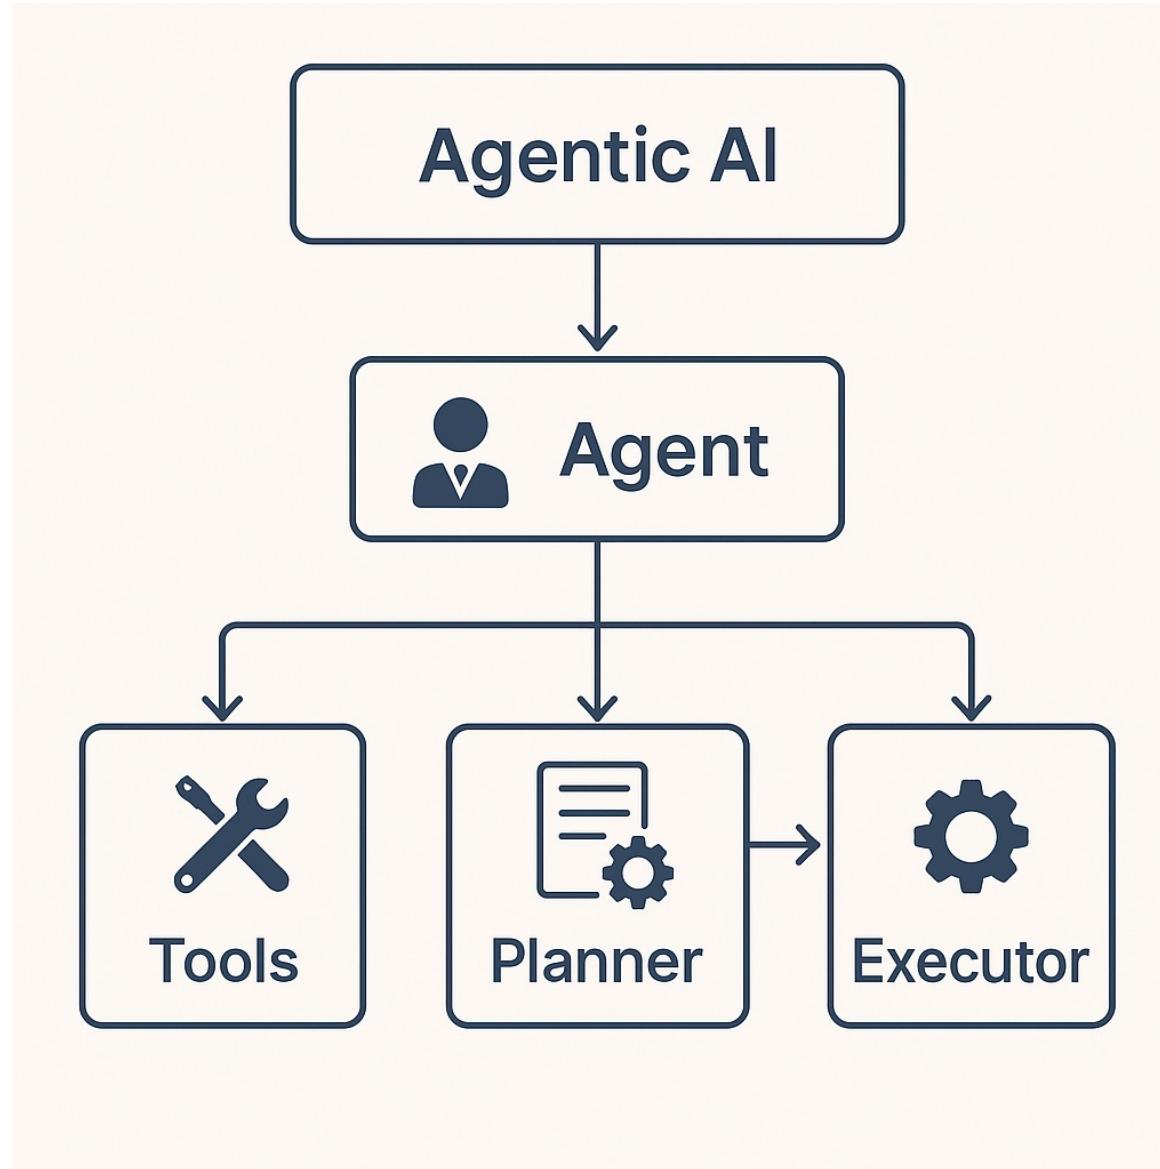
Core Components of Agentic AI

1. Agents

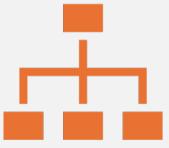
2. Tools

3. Planner

4. Executor



1. Agents



Individual units with specific roles or goals.



Can reason, learn, and make decisions



based on their goals.

2. Tools



Resources or capabilities available to agents.



For example: APIs, databases,



web search engines, or AI models.

3. Planner

Determines how to achieve the goal.

Decides which agents

should collaborate, and in what order.

4. Executor



Executes the
planned tasks.



Coordinates
interaction



between agents
and tools



to complete
tasks.

Core Components of Agentic AI

Component	Purpose
 Agent	An intelligent entity that receives instructions and decides what to do.
 Tool	A function or API the agent can use (like a search engine, database query, code interpreter, etc.)
 Planner	Breaks down a big problem into smaller tasks.
 Executor	Executes tasks one by one, using tools, and adjusts based on feedback.

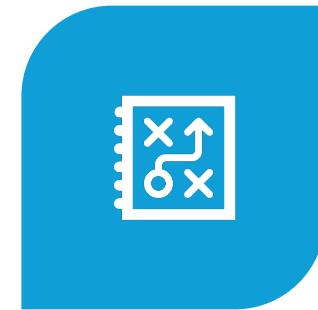
What Are Agent Design Patterns?



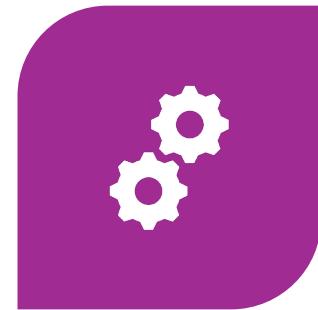
AGENT DESIGN
PATTERNS DEFINE



HOW AN AI AGENT



THINKS, PLANS,
AND



EXECUTES TASKS.

What Are Agent Design Patterns?

These patterns guide:

How the agent reasons

How tasks are broken down

How tools are used

How the agent adapts to feedback

Key Agent Design Patterns

ReAct (Reasoning and Acting)

Chain-of-Thought (CoT)

Plan-and-Execute (PnE)

ReAct **(Reasoning** **and Acting)**

The agent **reasons**

step-by-step and then

decides an action,

repeating until

the final answer is found.

Pattern



Thought → Action → Observation →



Thought → Action → ... → Final Answer

Chain-of-Thought (CoT)

The agent is encouraged
to explain its reasoning process
before giving an answer.
It doesn't act with tools
but reasons clearly.

What is Chain-of-Thought (CoT)?

Instead of jumping to the final answer,

the model is instructed to **think aloud**

by breaking the problem into

logical reasoning steps.

What is Chain-of-Thought (CoT)?

Improves accuracy, interpretability

often factual correctness,

especially for complex or

multi-step questions.

Plan-and- Execute (PnE)

The agent
first plans

the full
sequence of
tasks,

then
executes
each step.

What is Plan-and-Execute (PnE)?

Powerful agent pattern that:

First plans a sequence of subtasks

based on the user request.

Then **executes each step**,

often using tools, APIs, or functions.

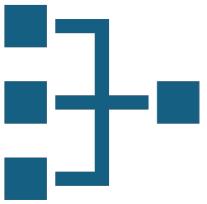
What is Plan-and-Execute (PnE)?

This improves

reliability,

transparency, and

modularity in AI task solving.



Why Orchestration Matters?

Connecting Multiple Agents and Workflows

Single LLM Calls to Systems of Agents



Move beyond a single prompt–response interaction



Need **orchestration**



The coordination of multiple agents



(each with a role) and tools



like databases, APIs, or



human-in-the-loop steps).

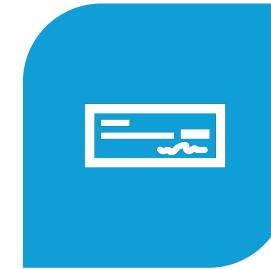
Healthcare RAG setup



RETRIEVER AGENT
PULLS POLICY DATA.



SUMMARIZER
AGENT DISTILLS IT.



COMPLIANCE
AGENT CHECKS
REGULATIONS.



AUDITOR AGENT
LOGS ACTIVITY.

Without orchestration

Agents act
independently

causing
duplication,

errors

untraceable
logic flow.

Core Benefits of Orchestration

Aspect	What It Enables	Example
Control	Structured task sequencing	Query → Retrieve → Reason → Respond
Transparency	Track state and data lineage	Audit trail of retrieved sources
Error Handling	Retry/fallback logic	Retry embedding store if vector DB fails

Core Benefits of Orchestration

Aspect	What It Enables	Example
Parallelism	Concurrent agent execution	Policy summarization and fact extraction in parallel
Scalability	Add/modify agents without breaking flow	Plug in a new “translator” agent easily

Analogy

Think of orchestration as a

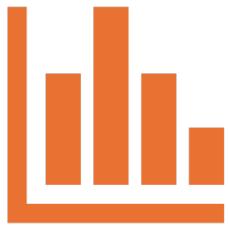
Conductor in an Orchestra

Agents (instruments) play different roles,

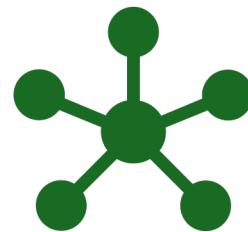
Conductor ensures

Timing, harmony, and handoffs.

LangGraph Architecture



State Graph



Nodes



Conditional Flows

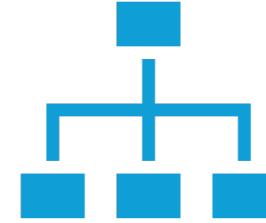
LangGraph Architecture



**Open-source
orchestration
framework**



for building Multi-agent
systems



**with dynamic control
flow.**

LangGraph Architecture

Graphs-as-programs

State transitions and

Node executions

Define how agents interact.

Core Components

Concept	Description	Analogy
Graph	The overall workflow connecting agents, tools, and logic paths	Blueprint of a system
Nodes	Executable units (functions, LLM calls, agents)	Steps or roles
Edges	Transitions between nodes (with conditions)	Arrows showing flow

Core Components

Concept	Description	Analogy
State	Shared memory or context that flows between nodes	Shared “workspace”
Conditional Flow	Decision logic that determines which path to take next	“If answer not found → fallback to external search”

Why LangGraph Is Powerful?

Combines

**LangChain
Runnables**

with

**graph-based
orchestration**

Why LangGraph Is Powerful?



Allows **visualization**



Debugging



Multi-step LLM



Reasoning

Why LangGraph Is Powerful?

Supports

Persistent
state

across
nodes

Why LangGraph Is Powerful?

Works with

async,

parallel,

conditional

executions

Why LangGraph Is Powerful?

Ideal for

Multi-agent collaboration

Retrieval workflows

Governance tracking

Visualization of Flow

[START]



[Retriever Node]

—> (if docs found?)

—> [Summarizer Node]

—> [Answer Node] —> [END]

 └—> [Fallback Search Node]

Orchestration + LangGraph = Cohesive Multi-Agent Systems

Orchestration provides

“*why and when*” (strategy).

LangGraph provides

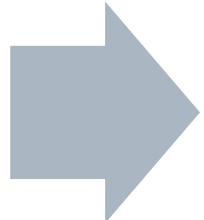
“*how*” (implementation pattern).

Concept Reinforcement

Concept	Implementation
Orchestration	Graph coordinates agents via state and edges
Multi-Agent Workflow	Retriever → Summarizer → Answer → Auditor
Conditional Flow	Branches if no docs retrieved
Governance & Audit	State tracks every step's activity
Extensibility	Add “Reviewer” or “Compliance Agent” easily

Multi-agent orchestration

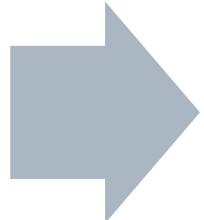
Instead of one
“do-everything”
agent,



**Use specialized
agents**
that collaborate

Multi-agent orchestration

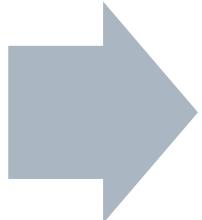
Planner:
decomposes goals
into steps



Researcher:
searches/reads
knowledge base &
tools

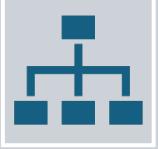
Multi-agent orchestration

Coder:
generates/fixes
code & tests



Reviewer:
checks
quality/safety

Multi-agent orchestration



Supervisor:
routes work among agents,



decides “done/hand back to user”

Multi-agent orchestration Benefit

Separation of concerns,

Easier testing/guardrails,

Better reuse.

Why LangGraph for orchestration?

State graphs with
nodes

(functions/LLMs/tools)

conditional edges

Checkpoints/memory

per session (resume,
replay)

Why LangGraph for orchestration?

Built-in **tool execution** nodes

Guardrails (insert moderation/validators as nodes)

Deterministic control flow you can test

LangSmith for Solution Architects



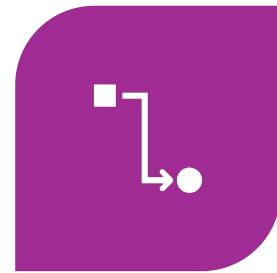
TRACE



EVERY LLM
CALL



TOOL STEP



IN A PIPELINE

LangSmith for Solution Architects



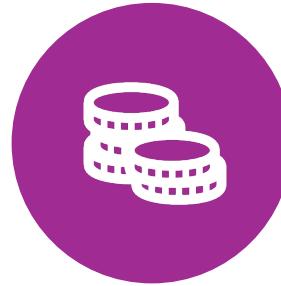
Debug



Failures and
regressions



with exact
inputs/outputs and



token/cost telemetry

LangSmith for Solution Architects



OBSERVE



LATENCY, COST,
AND



QUALITY OVER
TIME WITH

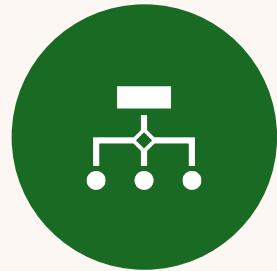


**DATASETS, EVALS,
AND FEEDBACK**

LangSmith for Solution Architects



GOVERN



RUNS WITH
**PROJECT-LEVEL
ISOLATION,**



REDACTION, AND



**AUDIT-READY
ARTIFACTS**

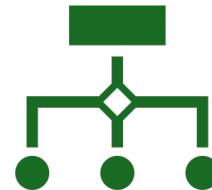


Why Observability Matters in Pharma/Healthcare

Audit & compliance



Preserve who ran what,



with which
prompt/model/parameters



when (key for SOPs, CAPAs,
validation).

Safety & privacy

Detect

Detect PHI leaks,

Apply

apply redaction,

**Tag and
quarantine**

tag and quarantine risky runs.

Reliability



Reproduce incidents with



full traces to speed up MTTR.

Cost control

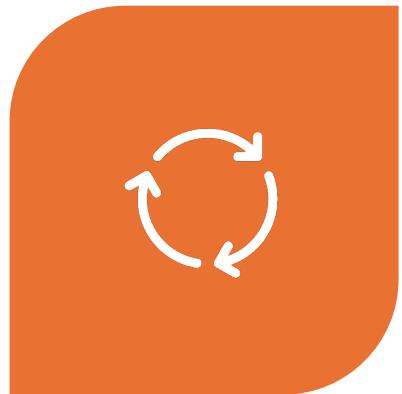


Track tokens + latency per component



(triage “where time/money goes”).

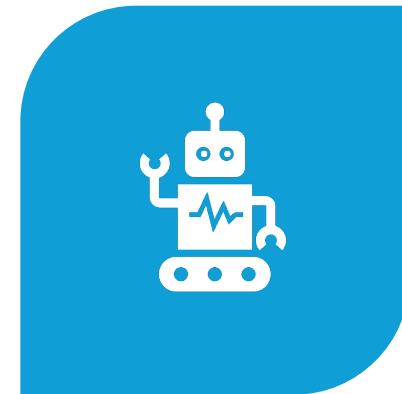
Quality



CLOSED-LOOP
IMPROVEMENTS



VIA USER FEEDBACK
+



AUTOMATED EVALS.

Core LangSmith Concepts (mental model)

Project



Logical workspace



(e.g., Viatris-Safety-QA).

Run / Trace

A single
execution
path

across LLM
calls,

tools,

retrievers,

chains.

Dataset

Versioned test

set of

input → expected output

Evaluation



AUTOMATED
GRADING



(LLM-AS-JUDGE OR
RULE-BASED)



OVER A DATASET.

Core LangSmith Concepts (mental model)

Feedback:

Human or
programmatic
signals

Categories,

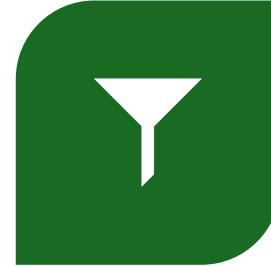
Numeric
ratings,

Comments

Core LangSmith Concepts (mental model)



TAGS/METADATA:



KEY-VALUE CONTEXT
TO FILTER,



SLICE, AND ANALYZE



(E.G., ENV=STAGING,
PHI=FALSE).

What is LangSmith?



Platform designed for



observability, evaluation, and deployment



of large-language-model (LLM)
applications and



agent workflows

Key aspects

Framework-agnostic.

LangSmith works

with any framework.

Tracing:

Capture each run of your LLM

Agent pipeline

inputs, outputs, metadata,

timing, cost

Debugging / root-cause

See

See where things go wrong

Tool

Tool calls, retries, errors

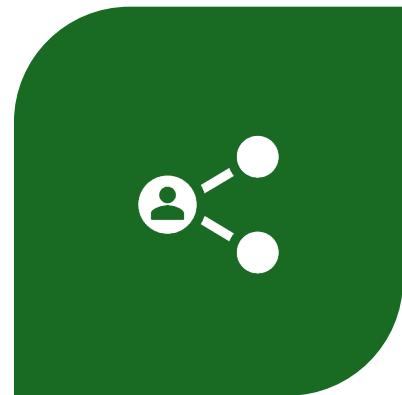
Enable

Enable alerts.

Evaluation + feedback



CREATE DATASETS OF INPUT →
EXPECTED OUTCOME,



EVALUATE YOUR LLM/AGENT
PERFORMANCE,

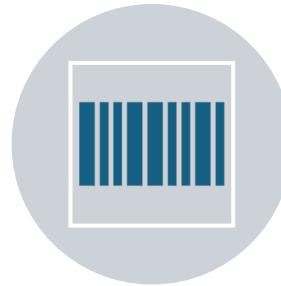


COLLECT FEEDBACK/HUMAN-
IN-LOOP.

Governance / auditability / deployment



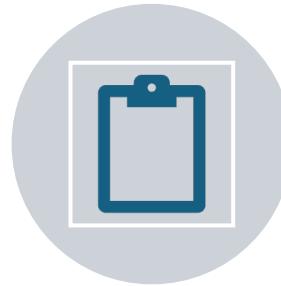
Support for enterprise-grade security,



Tag runs with metadata



Retention policies,



Redaction, export of traces for audits.

In short

If you build LLM-based workflows

Care about reliability, cost, compliance,

then LangSmith is the

observability + evaluation house around them.

LangChain vs LangGraph vs LangSmith

Tool	Primary Purpose	Best For	Key Strengths	Things to Watch/Limitations
LangChain	Framework for building LLM apps (chains, tools)	Linear or moderately complex LLM workflows (chatbots, RAG)	Easy to get started, rich integrations with models/data (Codebasics)	When flows get complex (branching, loops, state) it may become harder to manage
LangGraph	Graph-based orchestration of workflows/agents	Complex workflows with branching, loops, stateful agents	Explicit graph control, state management, agent support (LangChain)	Steeper learning curve; more infrastructure/complexity
LangSmith	Observability, tracing, evaluation, monitoring	Production deployments of LLM/agent workflows	Provides dashboards, tracing, feedback loops, audit-ready metrics (Medium)	It's not a workflow engine by itself; you still need to build the pipeline

Summary — Key Takeaways



USE LANGCHAIN
FOR BUILDING
YOUR



INITIAL LLM
WORKFLOWS



DOCUMENT Q&A,



SUMMARISATION,



RETRIEVAL

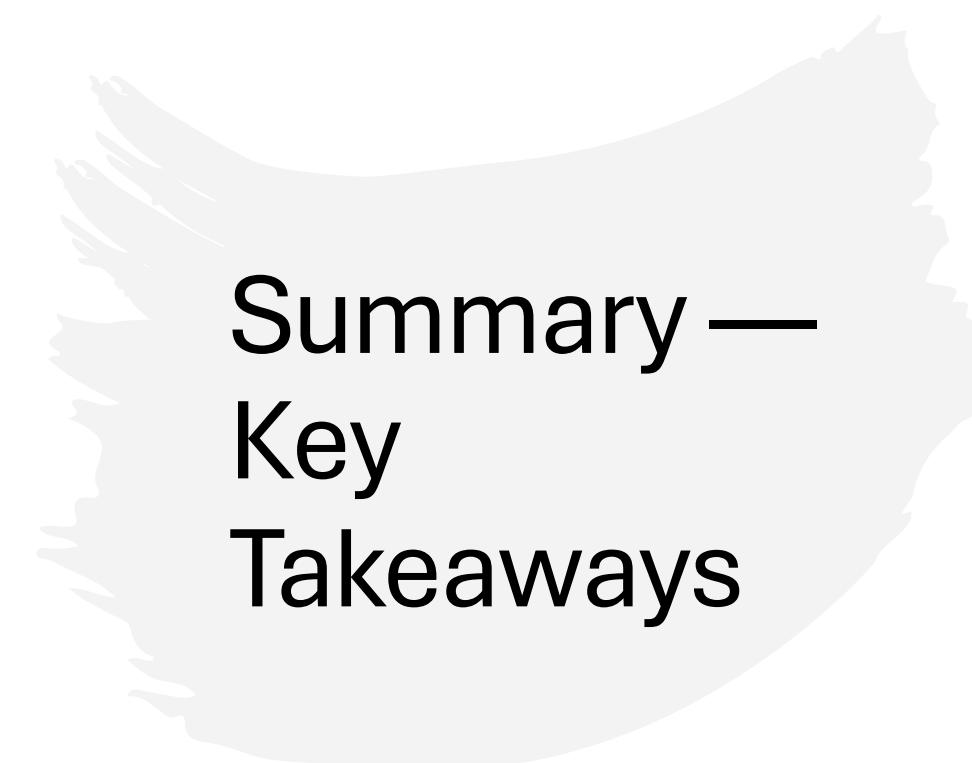
Summary — Key Takeaways

When workflows become multi-step,

stateful, or agentic

(for example PV triage agent with human review),

consider LangGraph for orchestration.



Summary — Key Takeaways

Regardless of workflow complexity,

embed LangSmith from early on

to enable tracing, cost/latency monitoring,

feedback loops, audits.

Summary — Key Takeaways

In the pharma/healthcare domain

you have additional demands

(compliance, PHI, audit trails)

which make observability, tagging,

review workflows essential

LangSmith is a strong fit here.

Happy Learning!!
Thanks for Your
Patience ☺

Surendra Panpaliya

GKTCS Innovations